

# lab 6

Lauren Waters

## All about functions in R

Functions are the way we get stuff done in R. We call a function to read data, compute stuff, plot stuff, ...

R makes writing functions accessible but we should always start by trying to get a working snippet of code first before we write our function.

## Today's lab

We will grade a whole class of student assignments. We will always try to start with a simplified version of the problem

```
# Example input vectors to start with
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

If we want the average we can use the `mean()` function

```
mean(student1)
```

```
[1] 98.75
```

Lets drop the lowest score. Can use the `min()` function to find the lowest value.

```
min(student1)
```

```
[1] 90
```

The `which.min` function may be useful here.

```
which.min(student1)
```

```
[1] 8
```

'-' will remove the value it's applied to (in this case the lowest score).

```
student1[-which.min(student1)]
```

```
[1] 100 100 100 100 100 100 100
```

Can put everything together to get the average grade (excluding the lowest value).

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Now, let's test on student 2.

```
mean(student2[-which.min(student2)])
```

```
[1] NA
```

Where is the problem? The same code doesn't work because they missed a homework, and there is a gap in the grades. By adding `na.rm=` argument, it strips all the NA from the sequence.

```
mean(student2, na.rm=TRUE)
```

```
[1] 91
```

```
mean(student3, na.rm=TRUE)
```

```
[1] 90
```

This isn't fair!

I want to stop typing out `student1`, `student2`, .... So, let's work with an input called 'x'

```
x <- student2
```

We want to overwrite the NA values with 0 - if you miss a homework, then you get 0% on the assignment.

Google and Chatgpt told me about the `is.na()` function.

```
x
```

```
[1] 100 NA 90 90 90 90 97 80
```

```
is.na(x)
```

```
[1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE
```

We can use logicals to index a vector. They pull out only the TRUE outputs

```
y <- 1:5  
y
```

```
[1] 1 2 3 4 5
```

```
y > 3
```

```
[1] FALSE FALSE FALSE TRUE TRUE
```

```
y[y > 3]
```

```
[1] 4 5
```

```
x[is.na(x)] <- 0  
x
```

```
[1] 100 0 90 90 90 90 97 80
```

This is my working snippet of code that solves the problem for all my example student inputs. Can replace x with each student (just reassign x everytime you want to change the student)

```
x <- student2
# Mask Na vlaues to zero
x[is.na(x)] <- 0
# Drop lowest score and get the mean
mean(x[-which.min(x)])
```

[1] 91

**Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped. Your final function should be adequately explained with code comments and be able to work on an example class gradebook such as this one in CSV format: <https://tinyurl.com/gradeinput> [3pts]**

```
grade <- function(x) {
# Mask Na vlaues to zero
x[is.na(x)] <- 0
# Drop lowest score and get the mean
mean(x[-which.min(x)])
}
```

Use this function:

```
grade(student1)
```

[1] 100

```
grade(student2)
```

[1] 91

```
grade(student3)
```

[1] 12.85714

We need to read the gradebook

```
gradebook <- read.csv("https://tinyurl.com/gradeinput", row.names =1)
gradebook
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	NA	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	NA
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	NA	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

Use `apply()` function to get the average grades for each student. For the second position (MARGIN), 1 = rows and 2 = columns.

```
ans <- apply(gradebook, 1, grade)
ans
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

**Q2. Using your `grade()` function and the supplied gradebook, Who is the top scoring student overall in the gradebook? [3pts]**

```
max(ans)
```

```
[1] 94.5
```

```
which.max(ans)
```

```
student-18  
18
```

**Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall? [2pts]**

Use the `mean()` and `apply` functions. If you use the `grade` function, then it will drop the lowest score.

```
mask <- gradebook  
  
mask[is.na(mask)] <- 0  
questions <- apply(mask, 2, mean)  
questions
```

```
hw1 hw2 hw3 hw4 hw5  
89.00 72.80 80.80 85.15 79.25
```

```
which.min(questions)
```

```
hw2  
2
```

We could also use the `sum()` function

```
questions <- apply(mask, 2, sum)  
questions
```

```
hw1 hw2 hw3 hw4 hw5
1780 1456 1616 1703 1585
```

```
which.min(questions)
```

```
hw2
2
```

**Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score (i.e. highest correlation with average grade score)? [1pt]**

You can add the `cor()` function

```
opt <- apply(mask, 2, cor, y=ans)
which.max(opt)
```

```
hw5
5
```